


# Understanding **Kubernetes** **Service Discovery & Ingresses**

Binura Gunasekara

SE @Platformer

Google Certified Cloud Architect

Kubernetes Certified Administrator




Kubernetes doesn't  
solve all the world's  
ills - @jbeda

# The **3 Commandments** of the k8s Network Model

- All Pods can communicate with all other Pods without using *network address translation* (NAT).
- All Nodes can communicate with all Pods without NAT.
- The IP that a Pod sees itself as is the same IP that others see it as.

Conceptually...

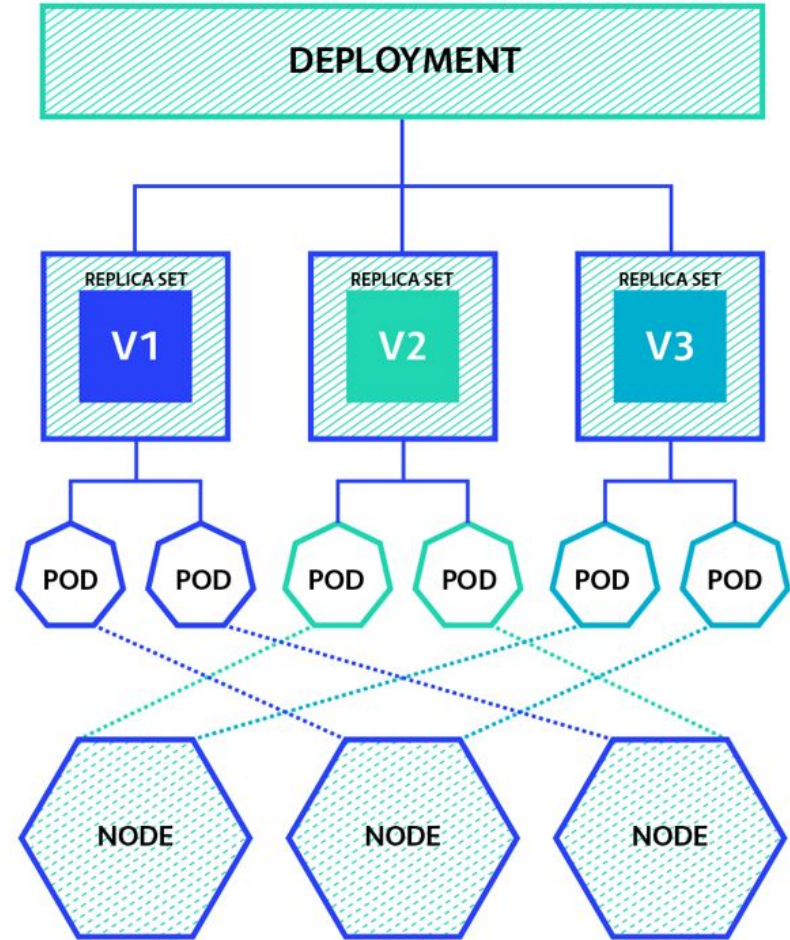


# Understanding Services

# Deployments

## Refresher

How do you we expose a number of running pods with a single point of entry?



## **ClusterIP:**

- Exposes the service on a cluster-internal IP. Choosing this value makes the service only reachable from within the cluster. This is the default

## **ServiceType.**

- Go-to Service type for internal communication between applications/microservices.

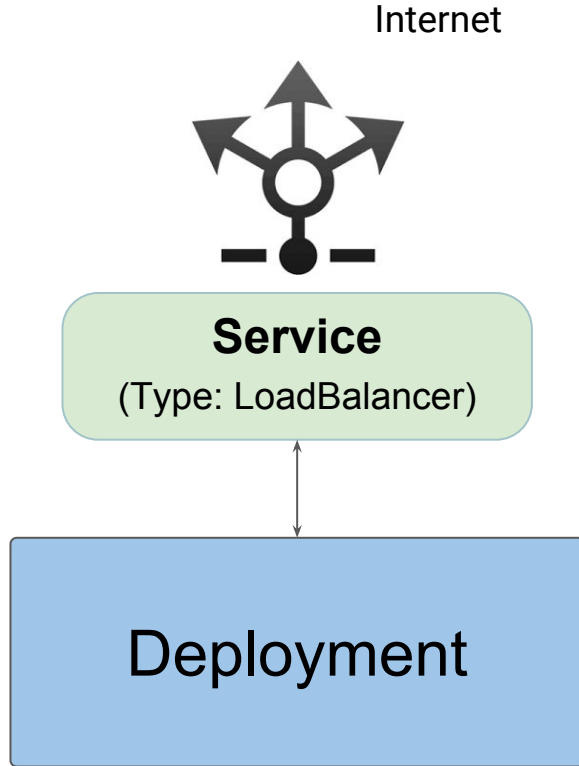
## NodePort

- Exposes the service on each Node's IP at a static port (the **NodePort**). A **ClusterIP** service, to which the **NodePort** service will route, is automatically created. You'll be able to contact the **NodePort** service, from outside the cluster, by requesting **<NodeIP>:<NodePort>**.
- Highly dependent on the Node's availability. Not recommended for Production



## LoadBalancer

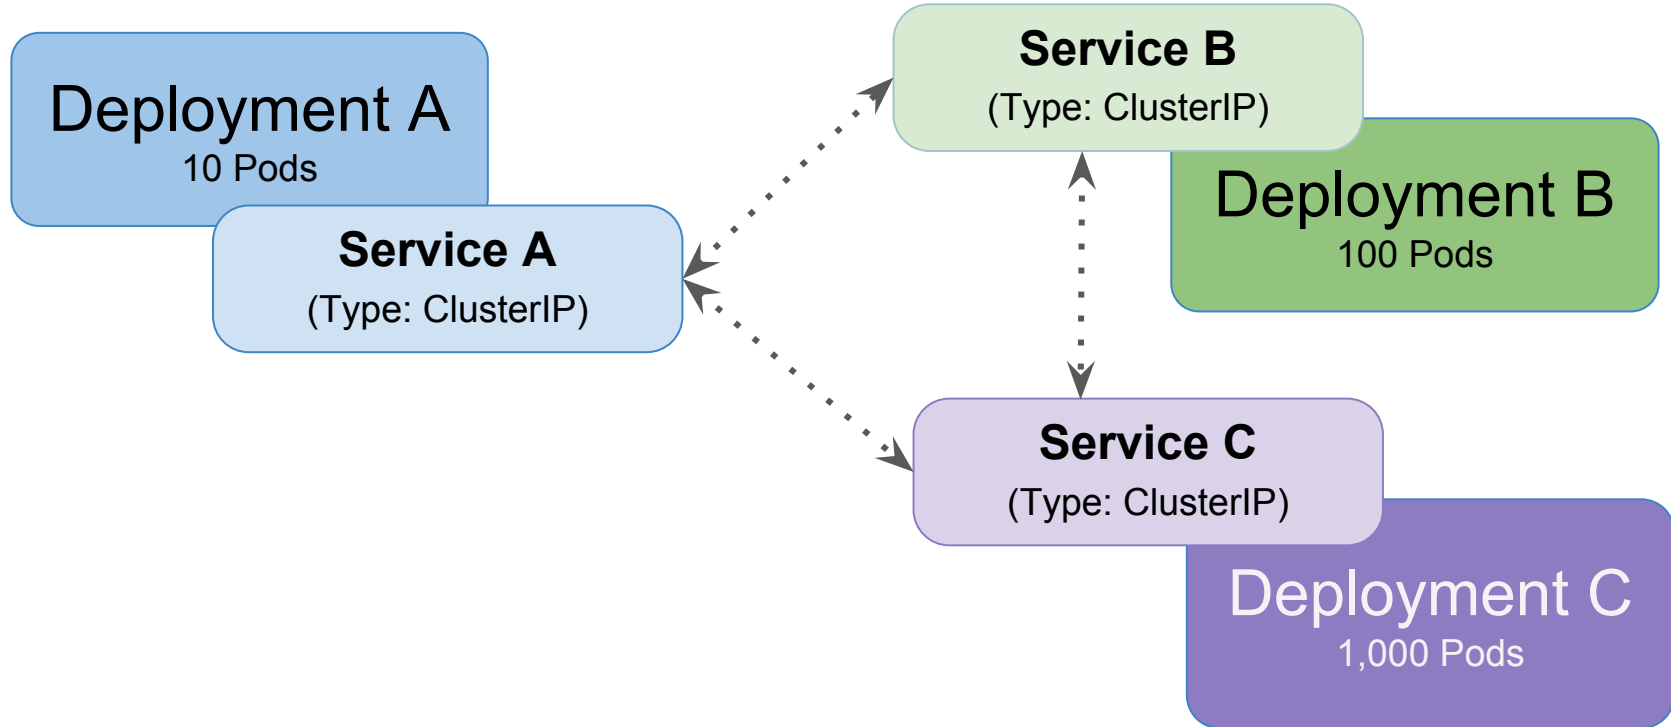
- Exposes the service externally using a **cloud provider's** load balancer. **NodePort** and **ClusterIP** services, to which the external load balancer will route, are automatically created.
- Requires a cloud provider. (There are on-prem solutions but they cannot be requested on demand with a LoadBalancer Type - requires a manual intervention)
- **Is it possible to share a LB between multiple applications? YES - Ingresses!**



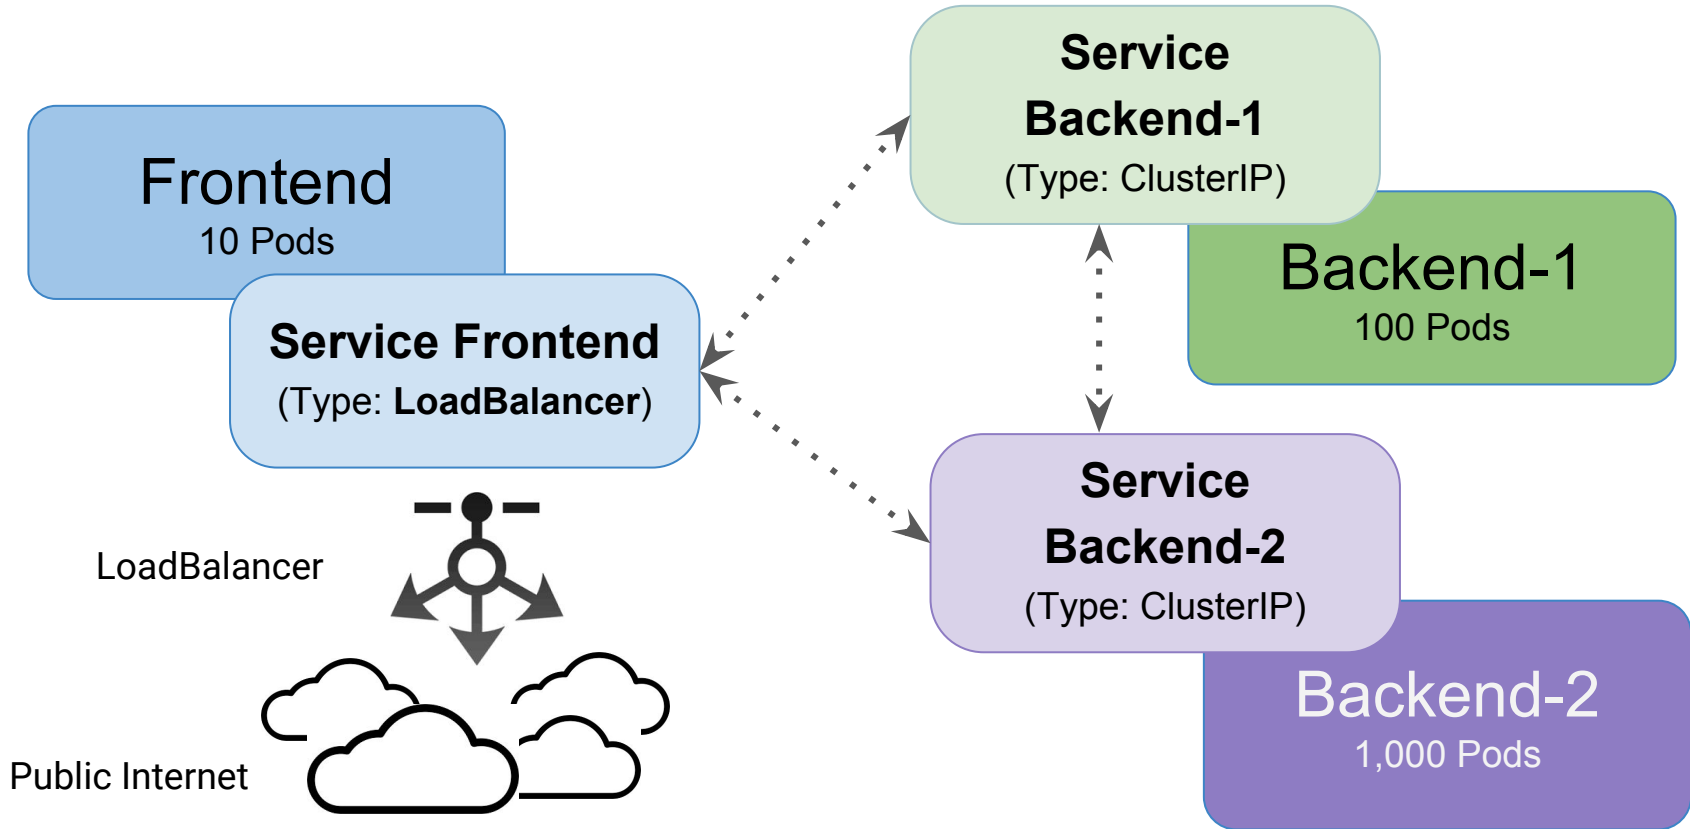
Manages the internal **NodePorts** and **ClusterIPs** as an *abstraction* -

Availability of Nodes is no longer a concern. Kubernetes does the mapping for us.

# ClusterIP - Internal Services



# Example Design



# Service Discovery

# There are **2** ways to discover Services

## DNS

- A Static DNS is created per each Service. (mapped to the ClusterIP by K8s internally).
- *Eg.* service-a can be reached by another Pod/Node using `http://service-a`
- If in a different namespace, `http://service-a.<namespace>`

## Environment Variables

- The Pod must be created after the service.
- The Kubelet adds a set of environment variables pointing to each Service ClusterIP to the Pod.

Eg.

```
SERVICE-A_PORT_80_TCP_ADDR=10.51.250.22  
SERVICE-A_PORT_80_TCP_PORT=80  
SERVICE-A_PORT_80_TCP_PROTO=tcp
```

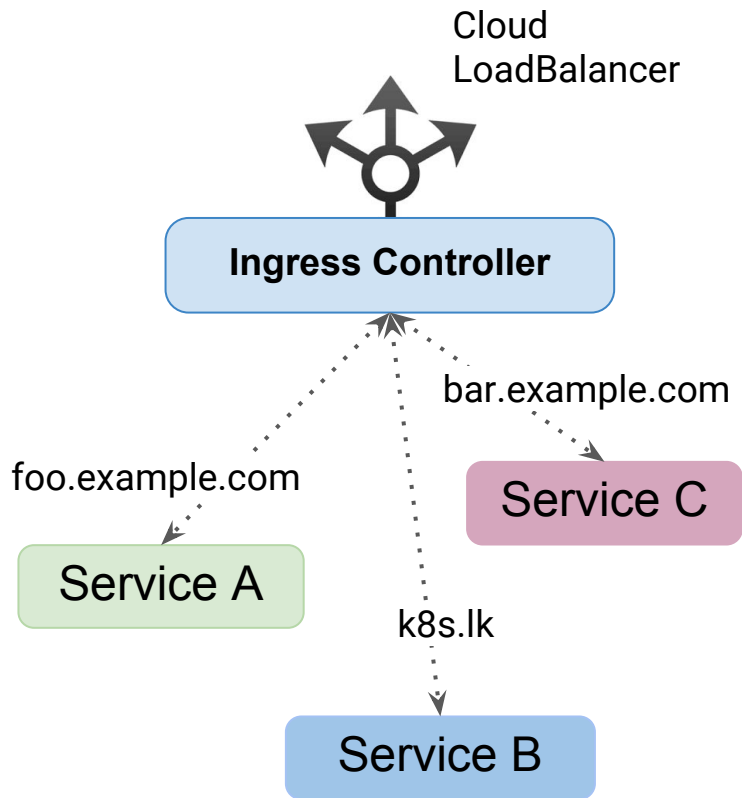
# Demo Time!

# Ingresses



# Ingesting Ingresses

- Ingresses expose HTTP/HTTPS routes from outside the cluster, to one or many services running inside the cluster.
- Popular use-cases:
  - Utilize a single Cloud Loadbalancer to serve external traffic to multiple applications.
  - Fan-out ingresses - quite useful when working with microservice-like architectures.
- Requires an **Ingress Controller** and a set of **Ingress Resources**.



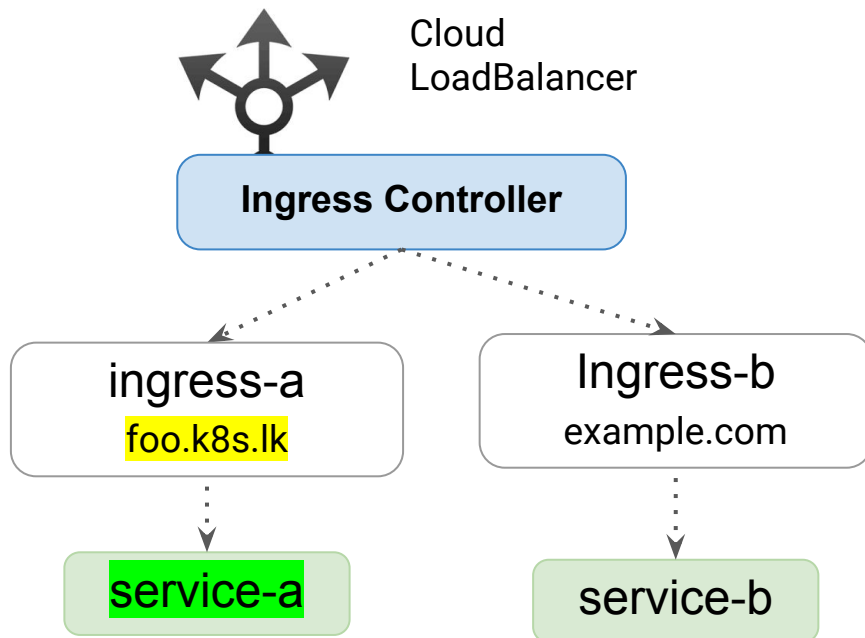
# The Ingress Controller

- Acts as a plugin Controller to the **kube-controller-manager**.
- Popular Ingress Controllers - *NGNIX, Traeffic, Istio Gateway* (The Nginx Ingress controller is different to the normal Nginx pods you've seen today.)
- Installing it is just a matter of **kubect! apply -f**.
- An Ingress Controller relies on a set of other resources called 'Ingresses' to figure out where to route which traffic.

# Single-service Ingress (Resource)

```

apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-a
spec:
  rules:
    - host: foo.k8s.lk
      http:
        paths:
          - path: /
            backend:
              serviceName: service-a
              servicePort: 80
  
```



# How do you point different DNS names at the Ingress Controller?

An Ingress Controller is assigned an **external IP** on creation. (This is in fact the Load Balancer's IP address - eg. 98.100.54.19).

Simply add the required CNAMEs to this IP Address. *And remember to make this IP static.*

CNAME RECORD	IP ADDRESS
foo.example.com	98.100.54.19
k8s.lk	98.100.54.19
bar.example.com	98.100.54.19

# Fan-out Ingress (Resource)

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: fanout-ingress-example
spec:
  rules:
    - host: shopping.lk
      http:
        paths:
```

```
    - path: /cart
```

```
      backend:
```

```
        serviceName: cart-service
```

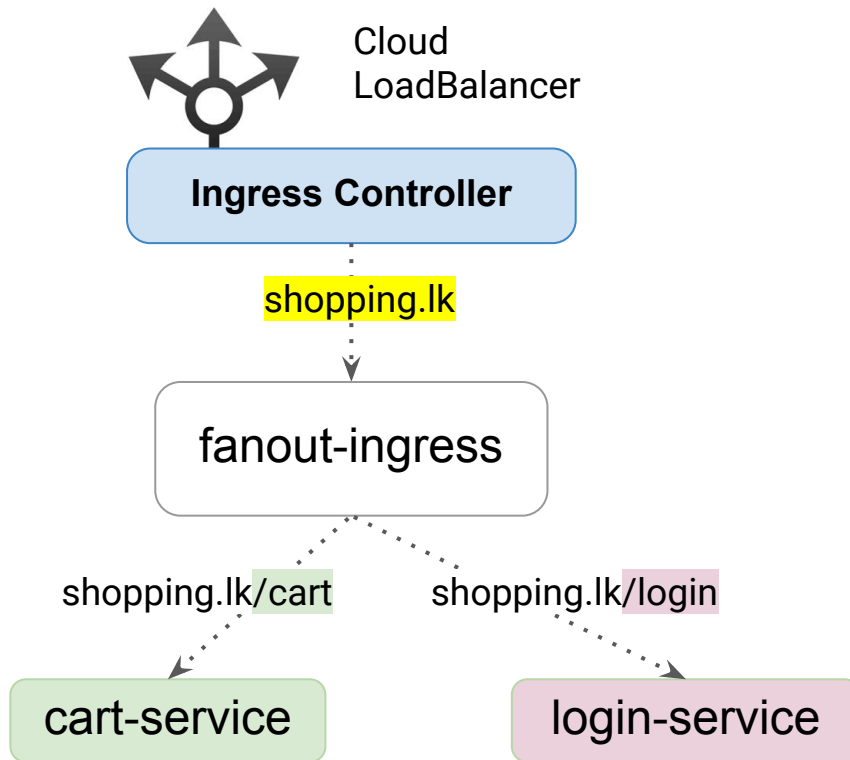
```
        servicePort: 5000
```

```
    - path: /login
```

```
      backend:
```

```
        serviceName: login-service
```

```
        servicePort: 8080
```



# Thanks!

- Get these slides from <https://github.com/BinuraG>
- Refer to the Docs at [kubernetes.io/docs](https://kubernetes.io/docs)

Contact us through [platformer.com](https://platformer.com) for consultation.

[binura.g@platformer.com](mailto:binura.g@platformer.com)